Parallel Fringe Search

Christian Zeman & Lukas Mosimann

ETH Zürich

Design of Parallel and High-Performance Computing

December 14, 2013

Overview

1. Algorithm
2. What we have done
3. Locking concepts
4. Benchmarks

Fringe Search

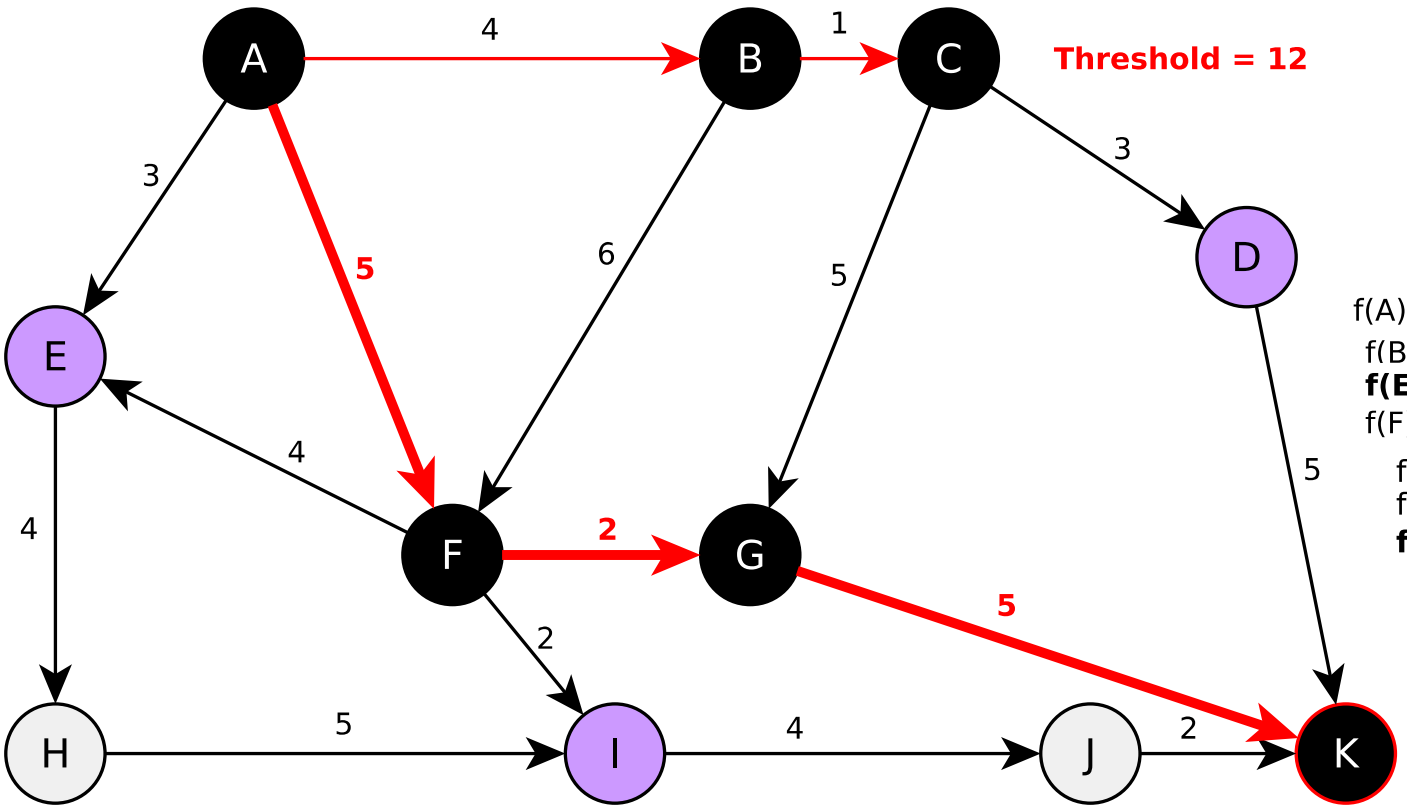

What we have done

- Serial implementation of fringe search (much faster than Boost A*)
- Parallel implementation with Open MP
  - 2 different locking concepts
  - Locks implemented using inline assembly (faster than Open MP locks)
- Benchmarking
  - Strong scaling
  - Weak scaling
  - Path quality

Locking concept: Deadlock prevention

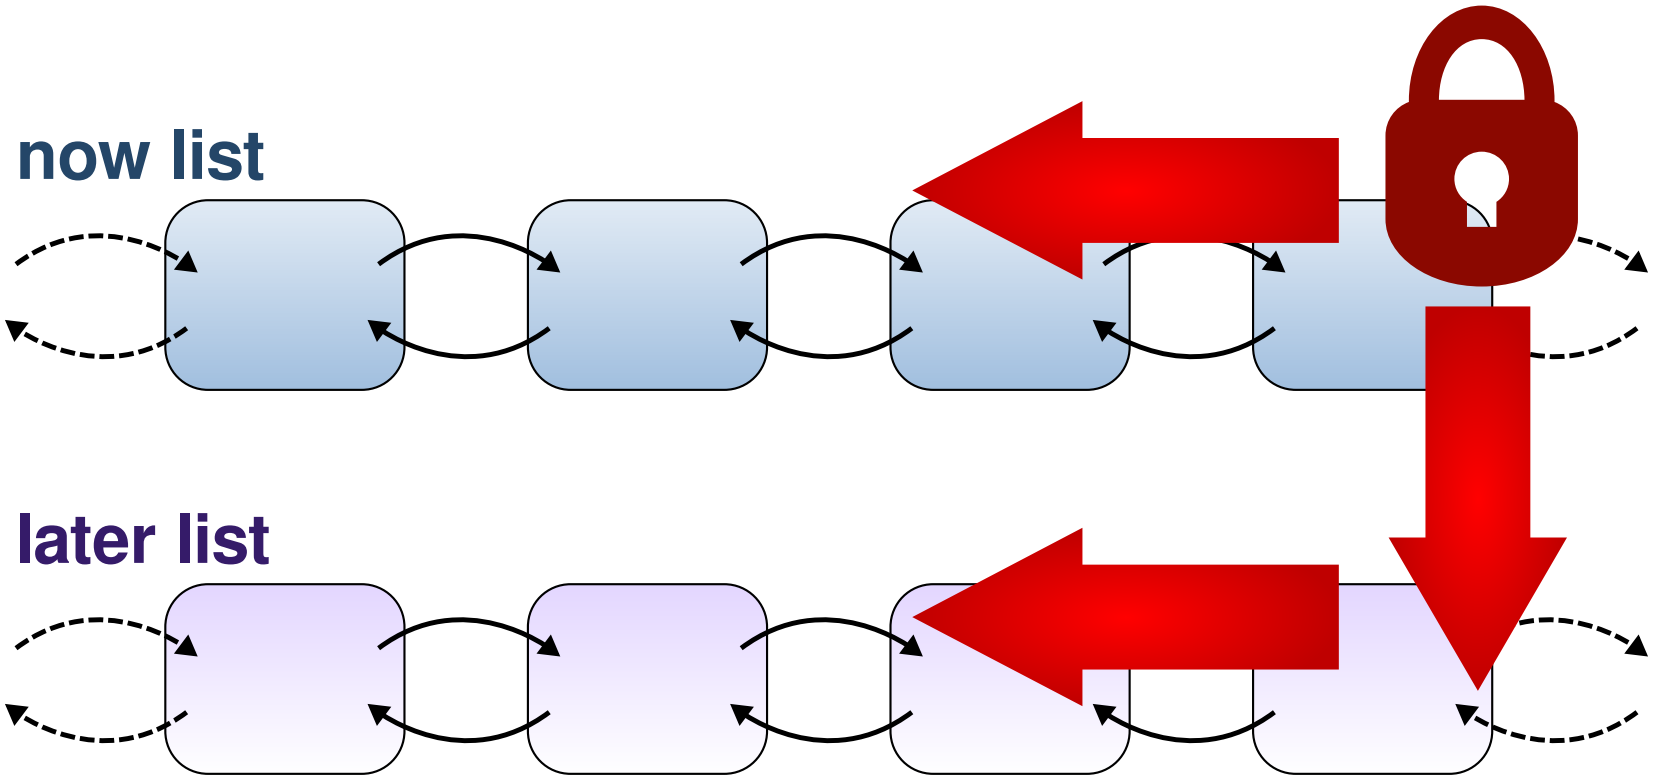

Locking concept: Insert node

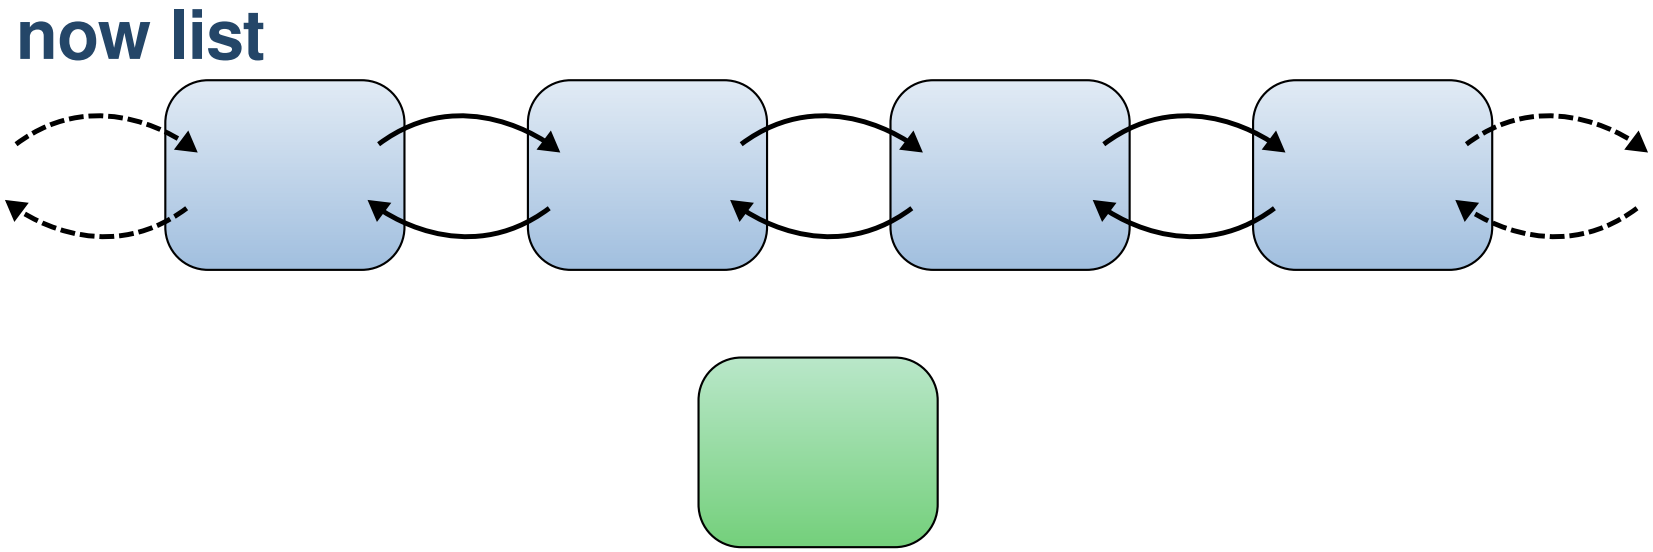

Locking concept: Insert node

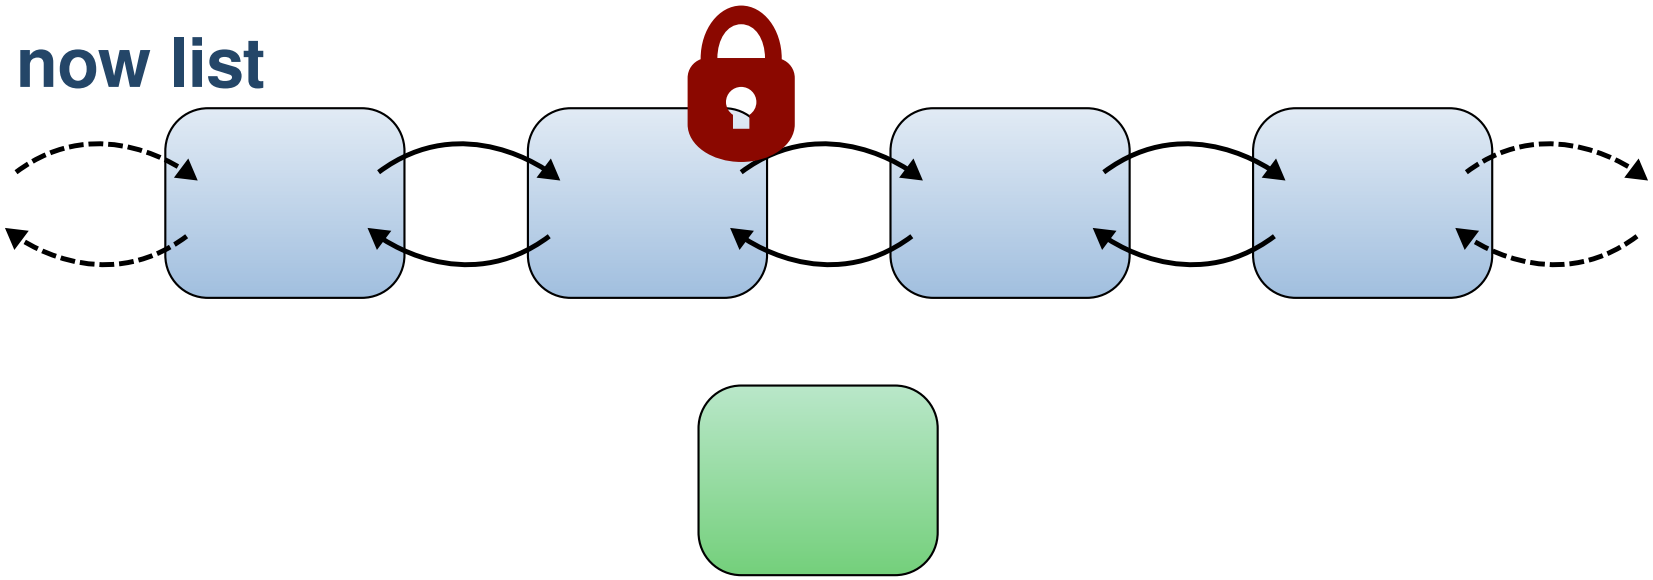

Locking concept: Insert node

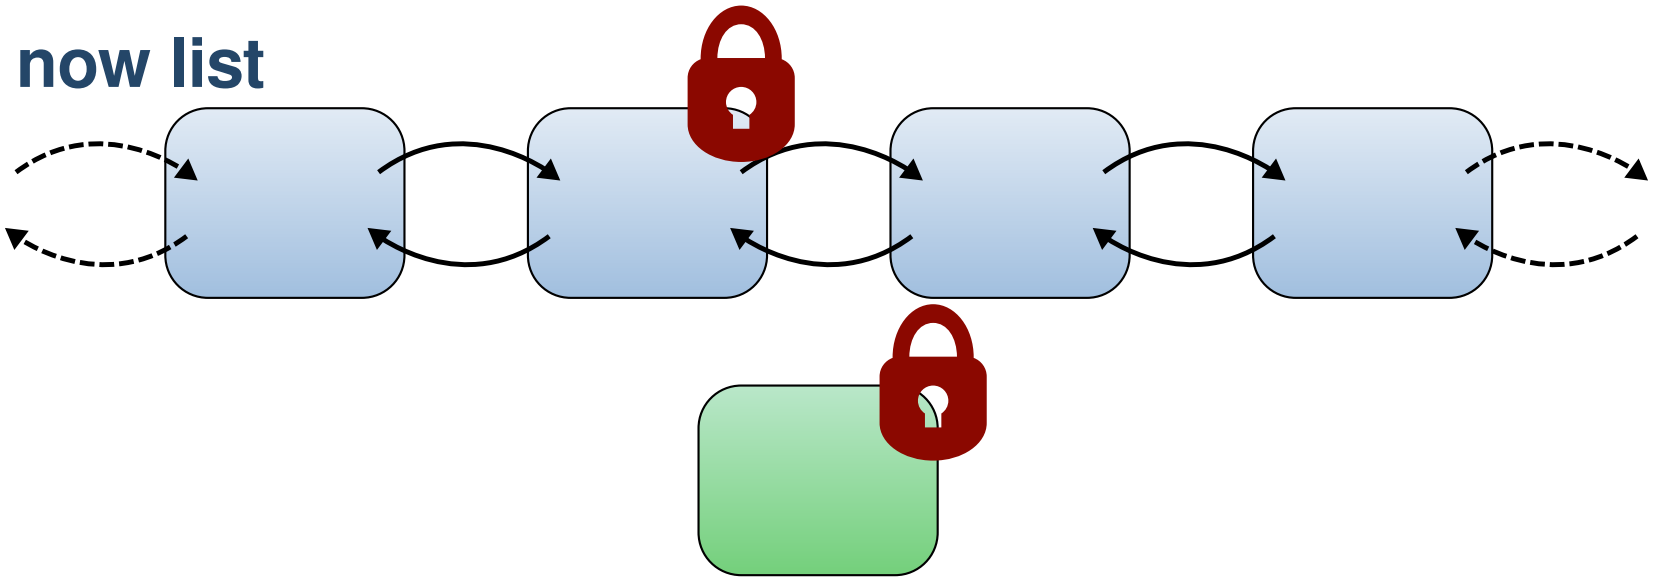

Locking concept: Insert node

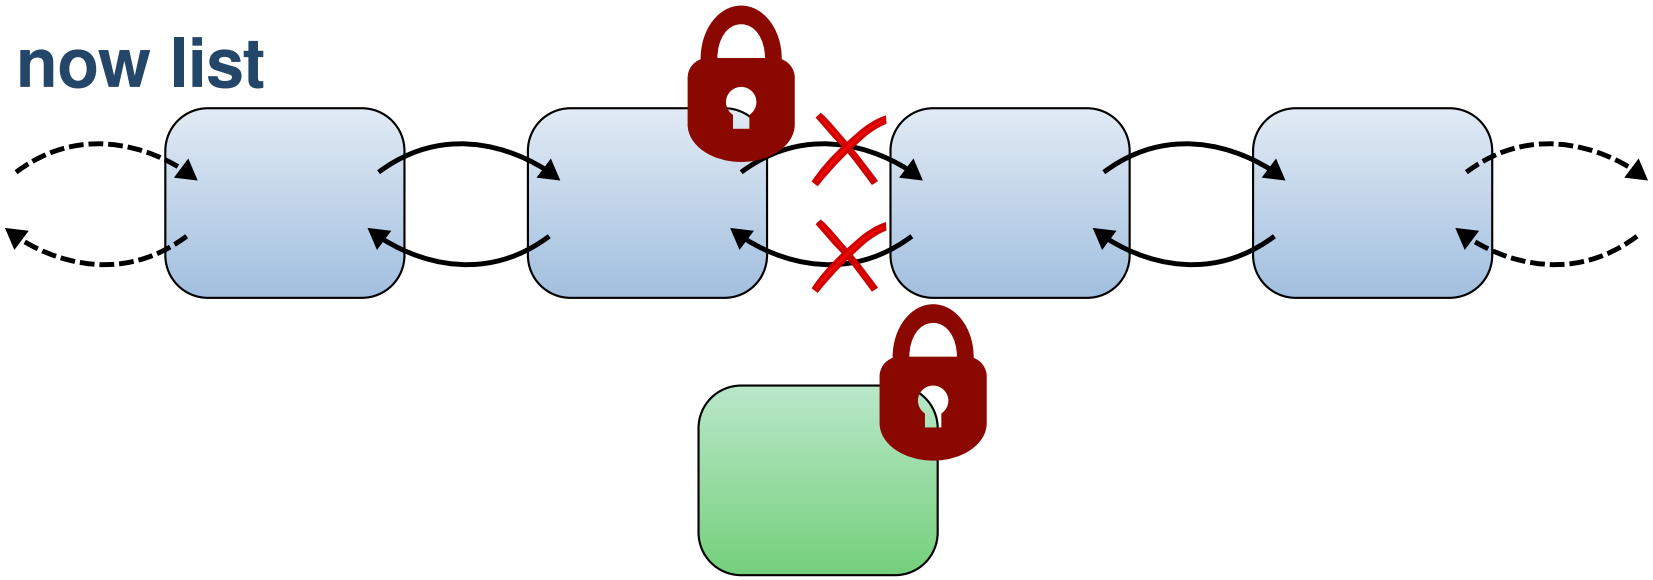

Locking concept: Insert node

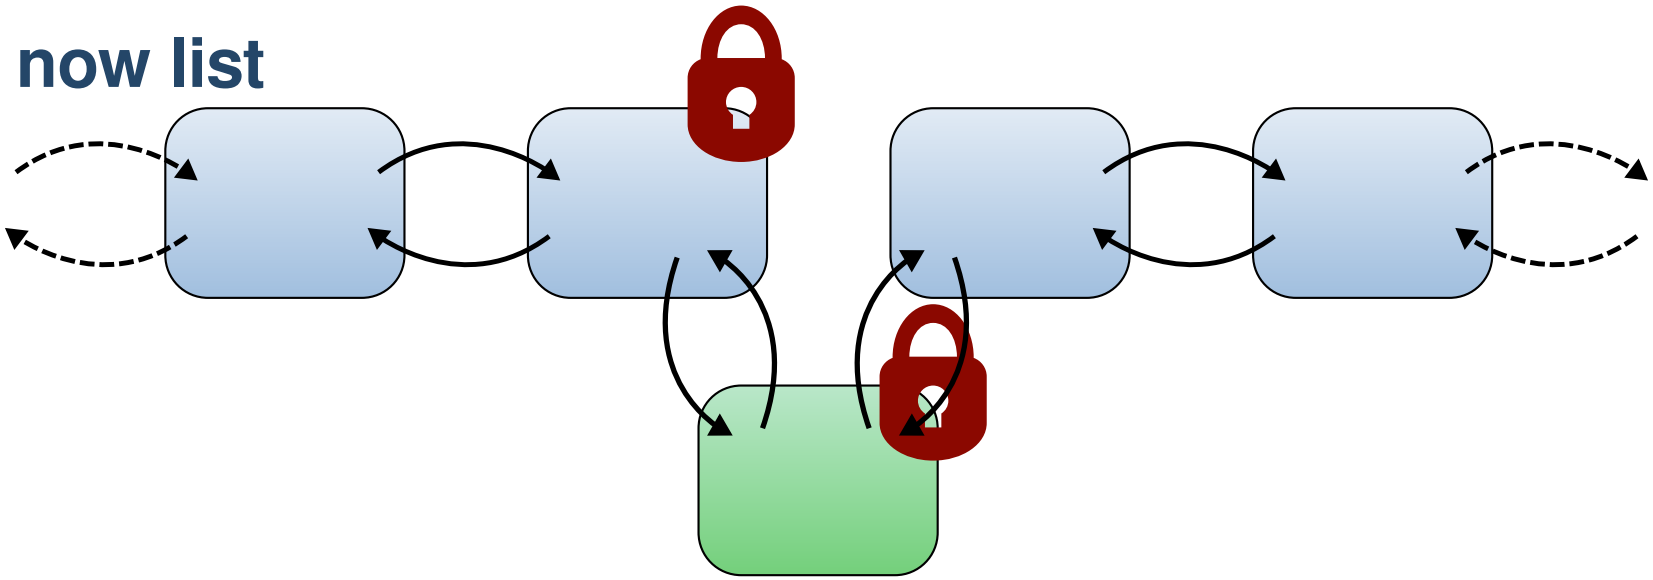

Locking concept: Insert node

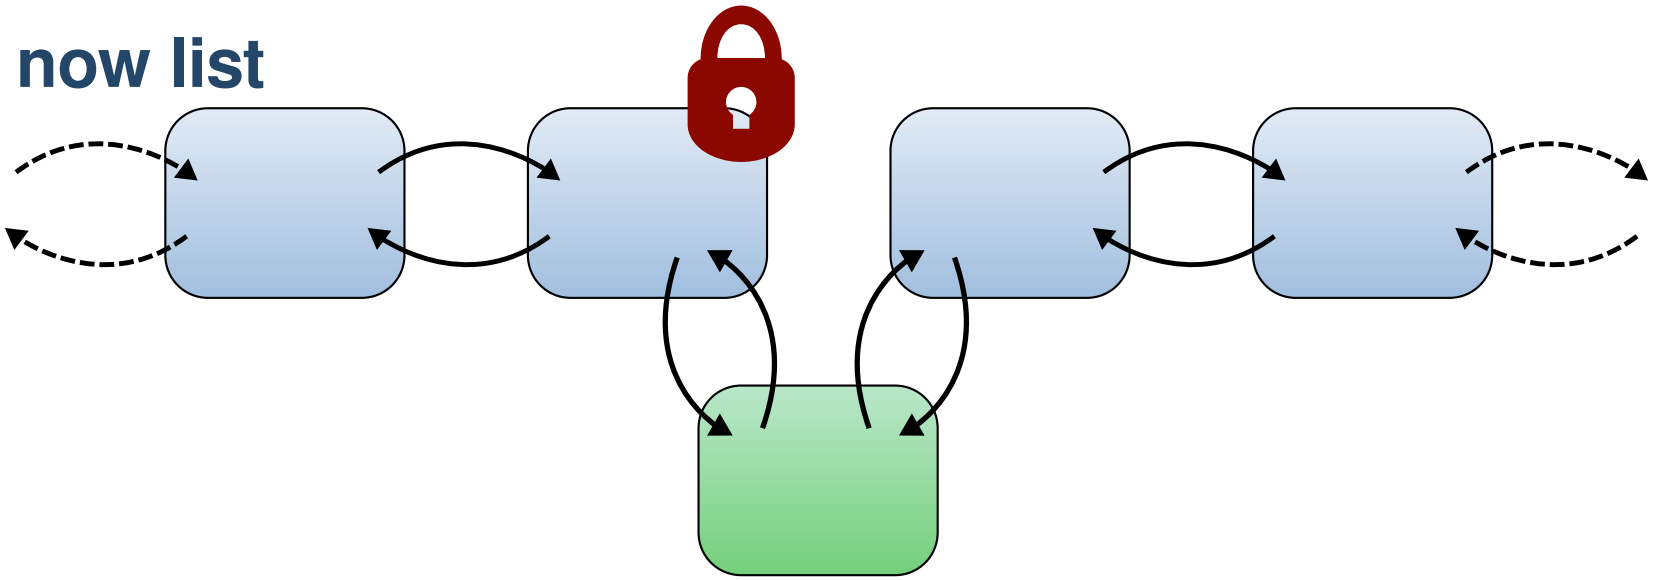

Locking concept: Insert node

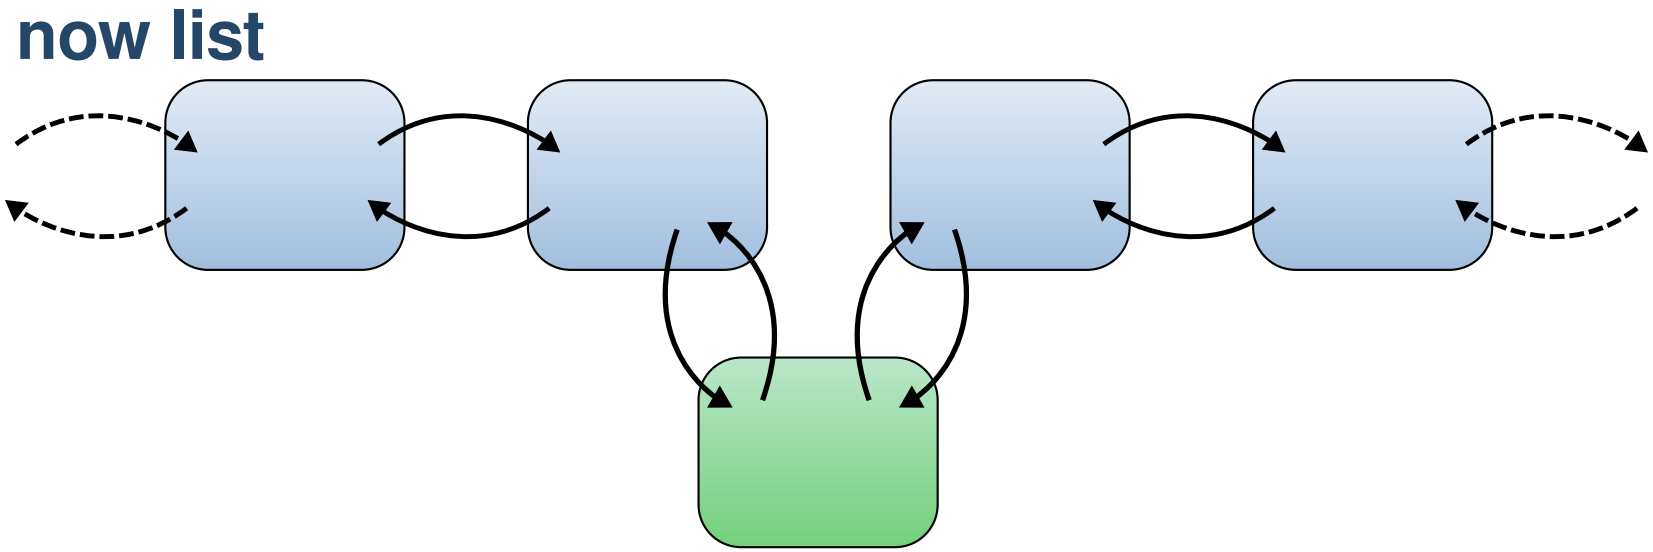

Locking concept: Remove node

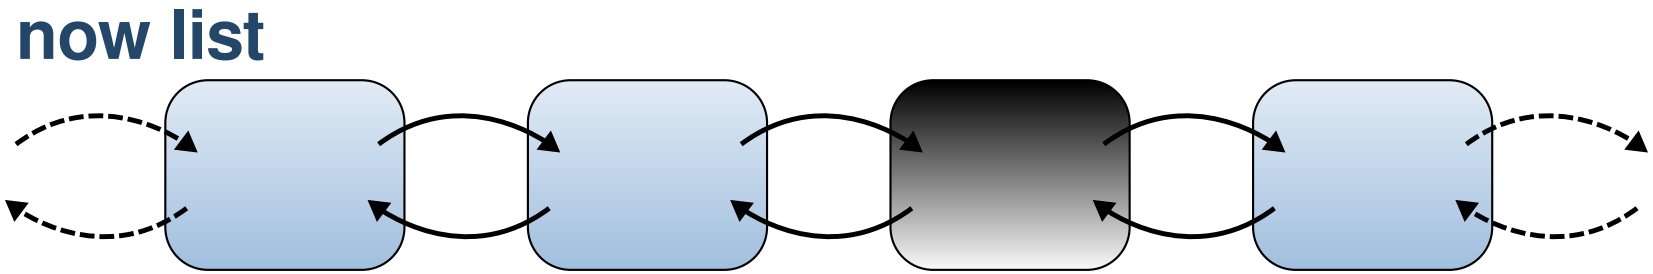

Locking concept: Remove node

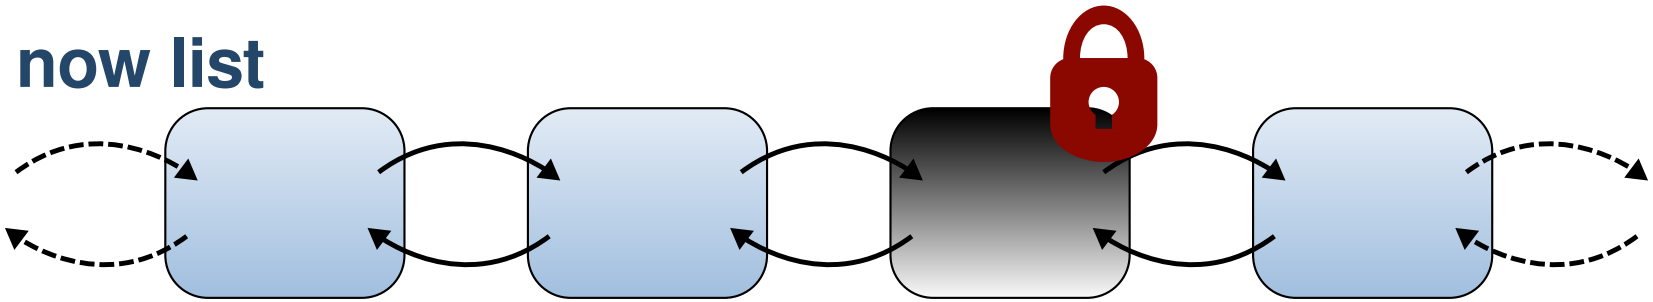

Locking concept: Remove node

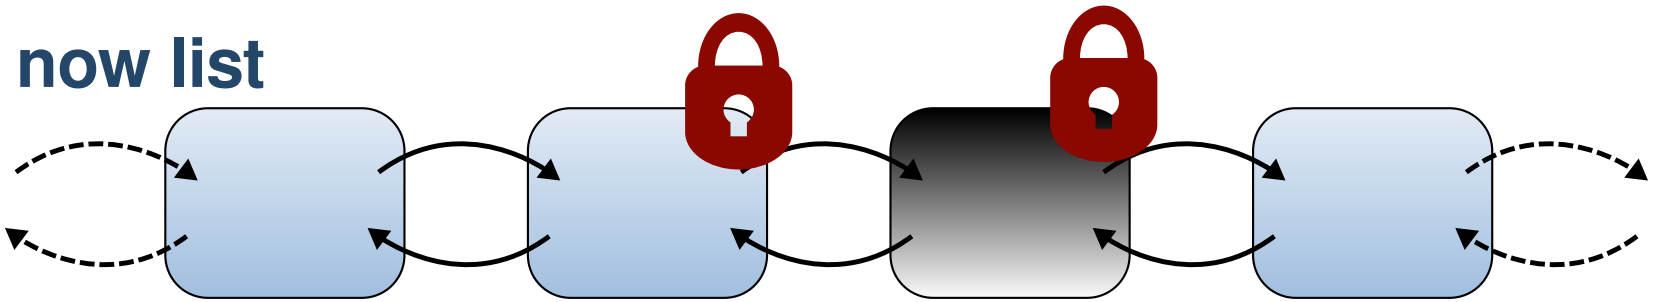

Locking concept: Remove node

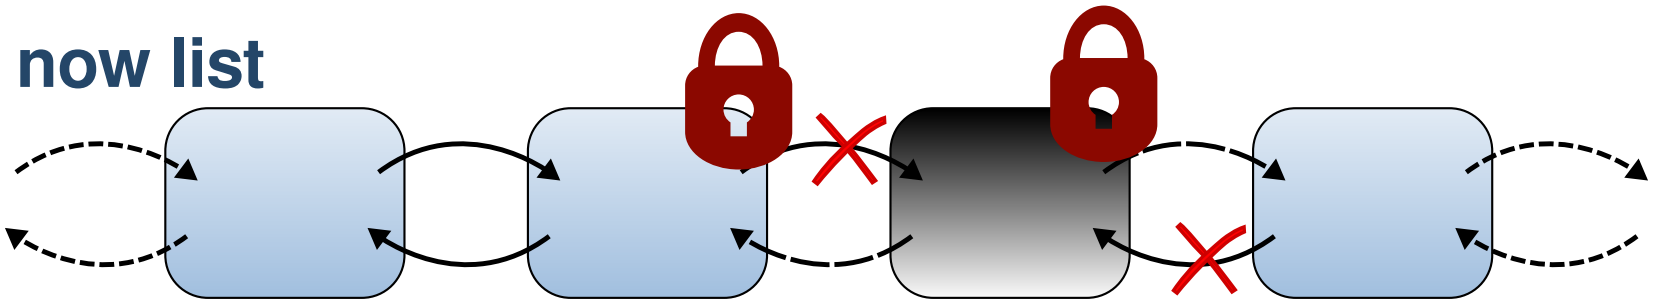

Locking concept: Remove node

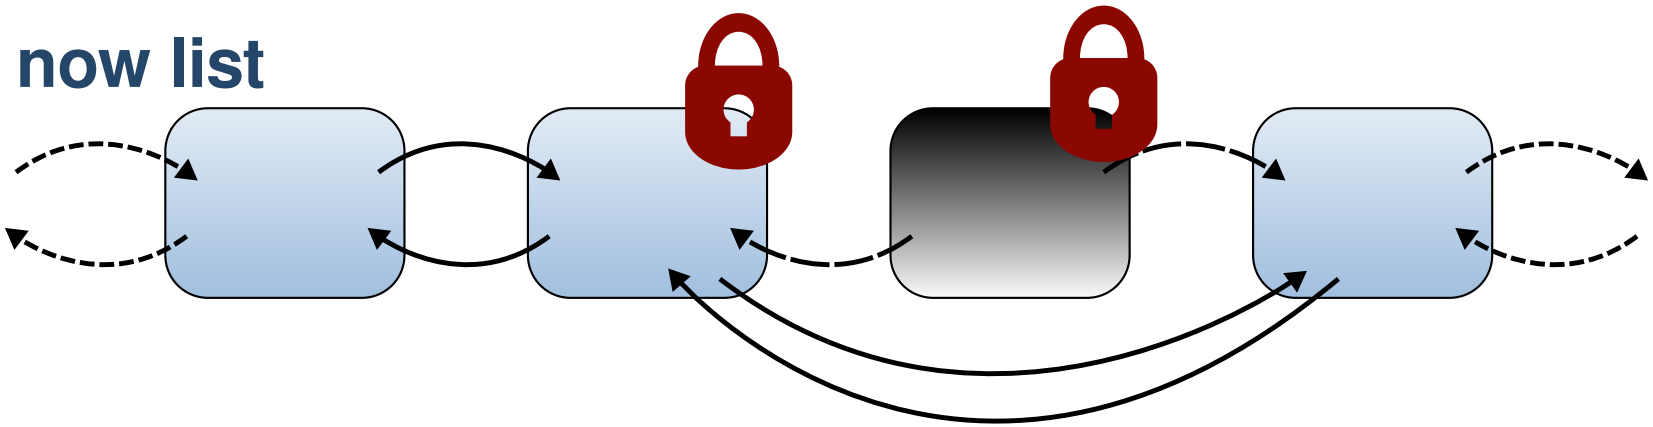

Locking concept: Remove node

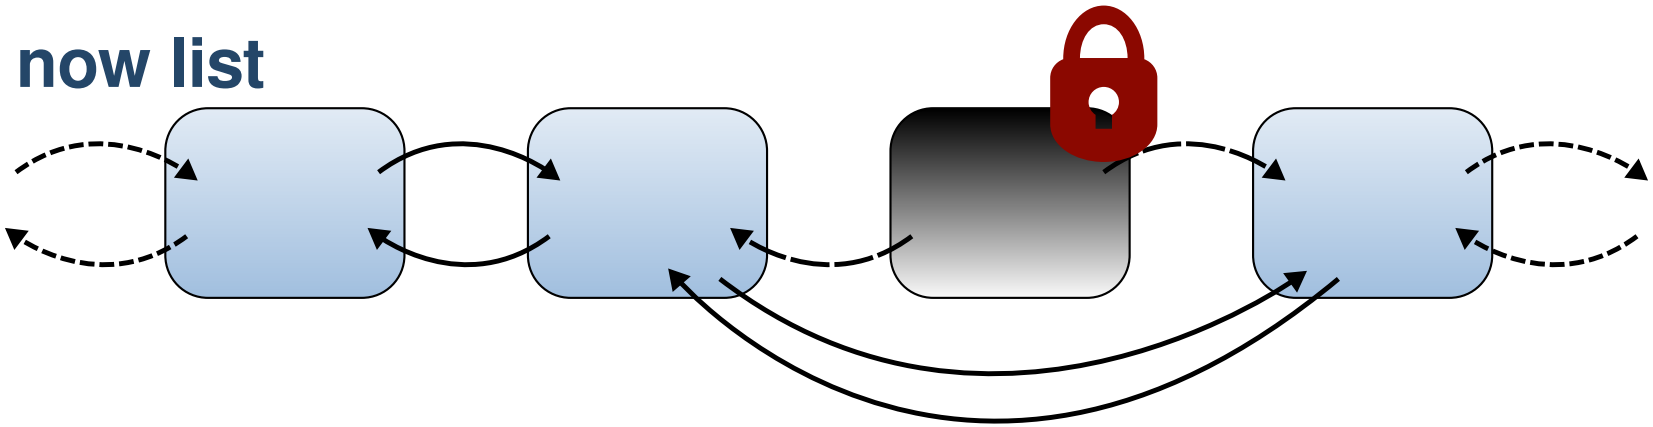

Locking concept: Remove node

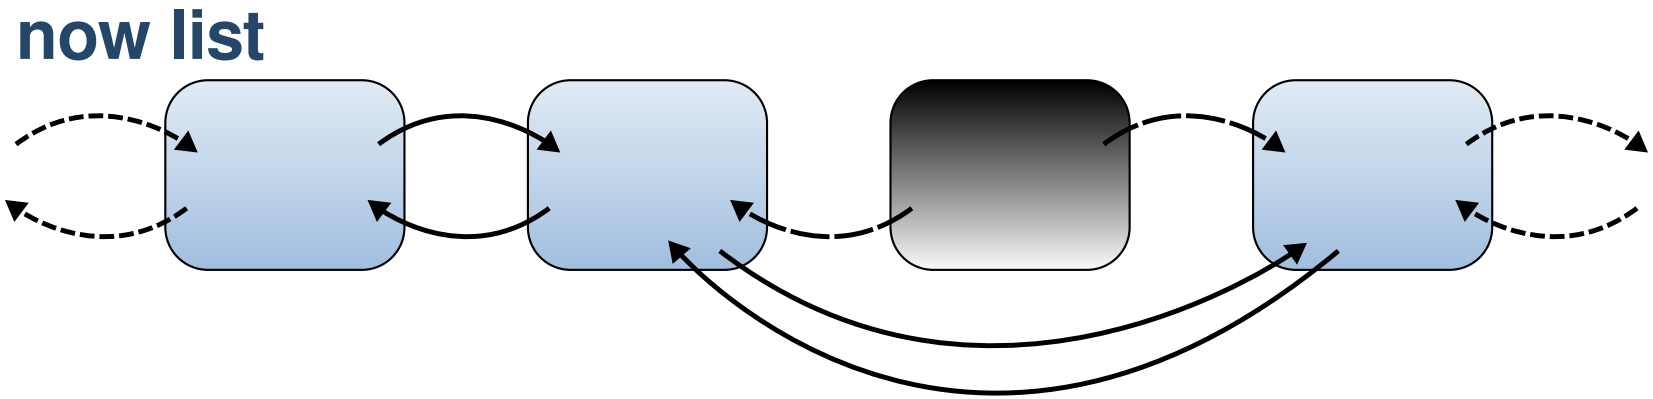

Threshold vs. Path Length

- Threshold = Threshold + **0.1**
- Threshold = Threshold + **1**
- Threshold = Threshold + **10**

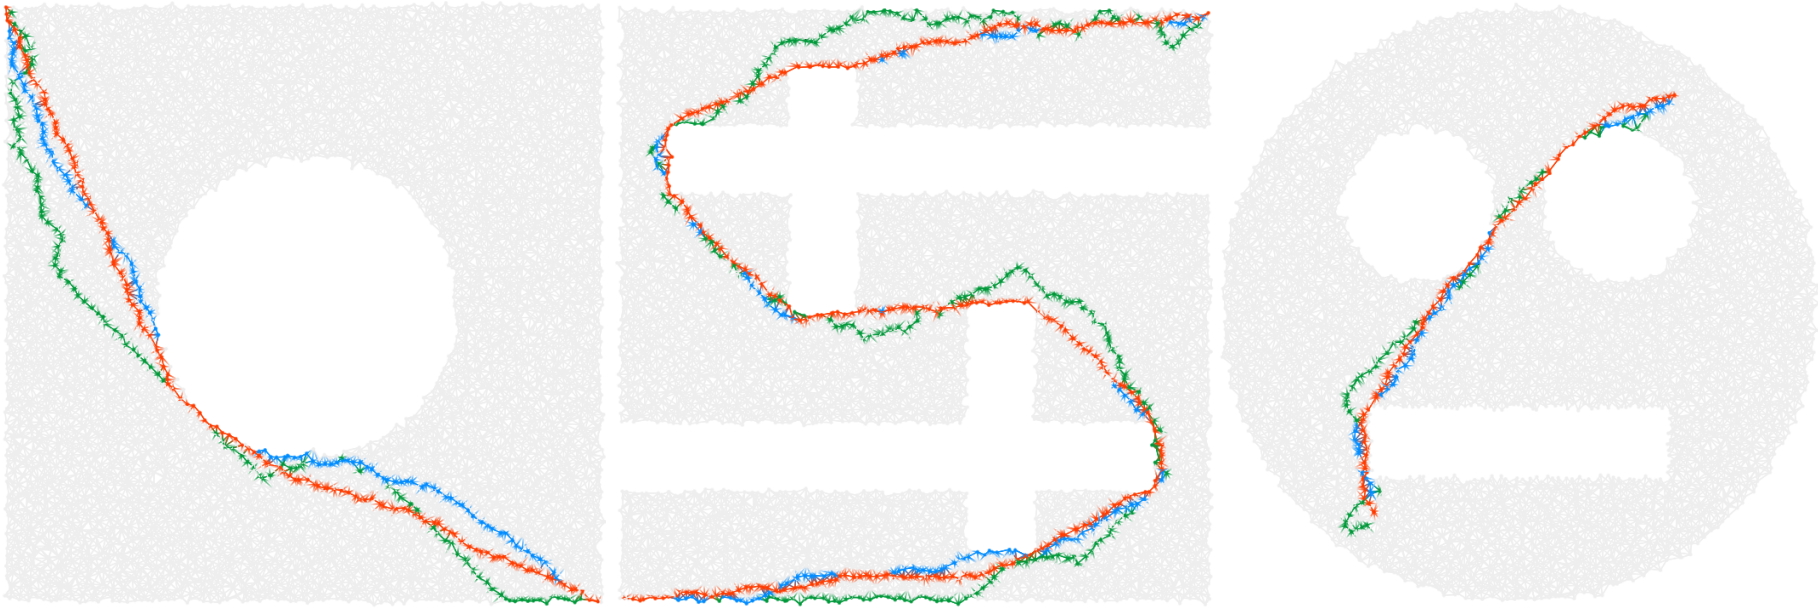

Benchmark runtime

- Runtime vs. A* from Boost Graph Library
- Using "cross graph" and 400 - 10'000 nodes
- Threshold += 1
- 20 runs per size
- Intel Core 2 P9600 @ 2.53 GHZ

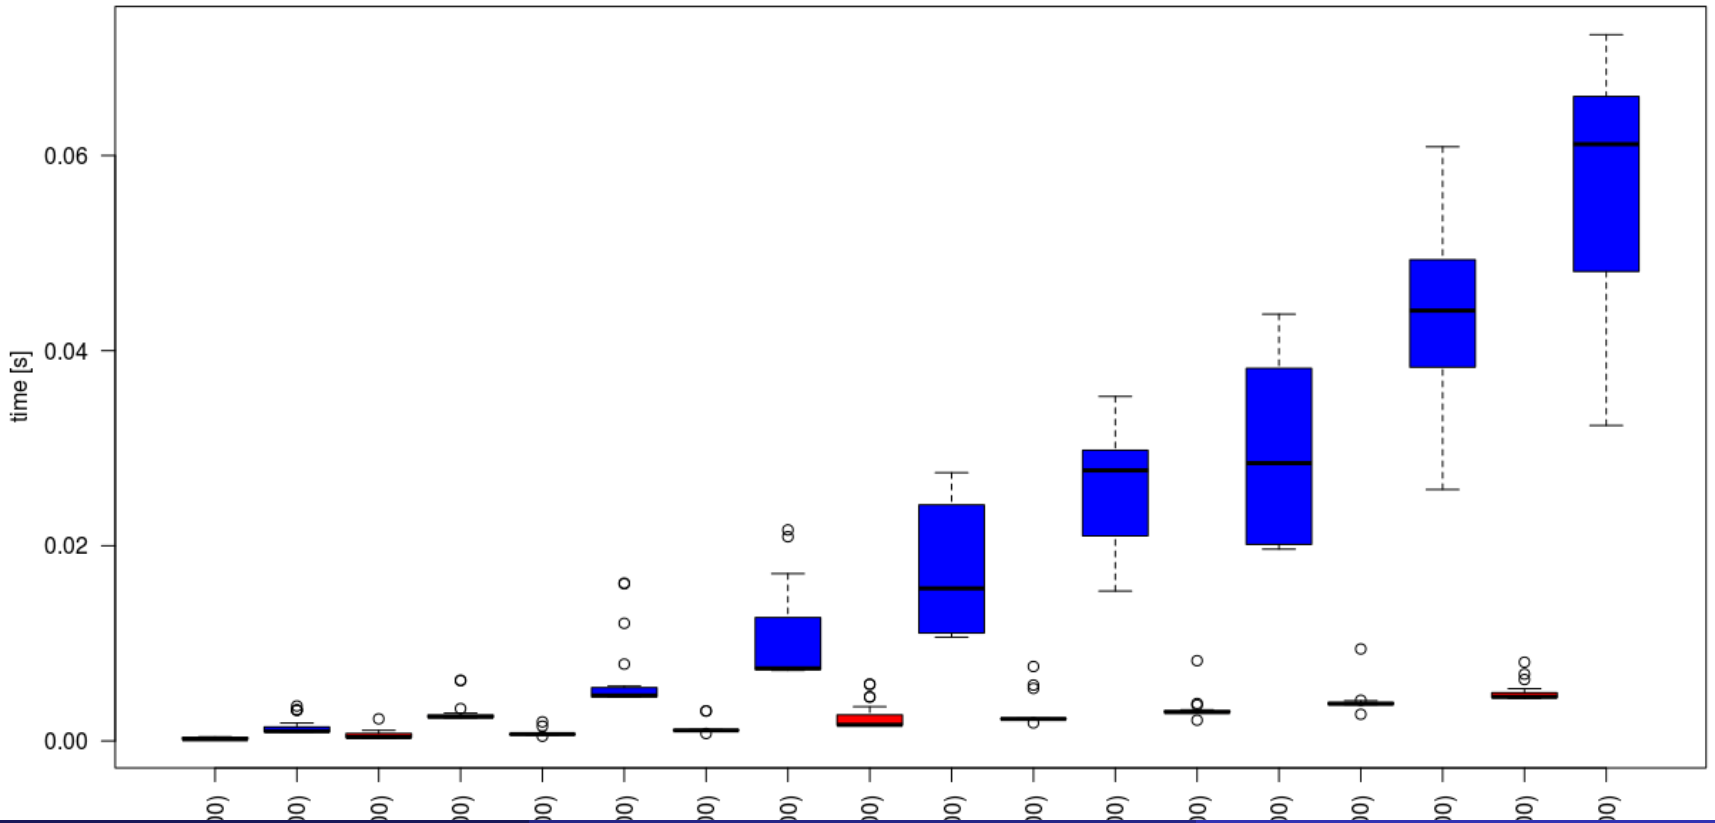

Benchmark path length

- Runtime vs. A* from Boost Graph Library
- Using "cross graph" and 400 - 10'000 nodes
- Threshold $+= 1$
- 20 runs per size
- Intel Core 2 P9600 @ 2.53 GHZ

Benchmark path length

# Nodes	A-Star	Fringe	Difference	Difference in %
400	69.731	70.6903	0.9593	1.37 %
900	101.071	102.638	1.567	1.55 %
1600	133.304	136.326	3.022	2.27 %
2500	165.47	167.891	2.421	1.46 %
3600	195.363	198.115	2.752	1.41 %
4900	227.957	231.554	3.597	1.58 %
6400	258.5	261.49	2.99	1.16 %
8100	288.961	293.483	4.522	1.56 %
10000	321.263	325.853	4.59	1.43 %

Project Status

What we still have to do:

- Implement parallel version
- Benchmarking (strong/weak scaling)
- Threshold fine-tuning
- Write report

The End